


It passed its own tests. *Then it got audited.*

A tool that plans software before it is written. It had its own checks and every one passed. A second reviewer that had never seen it built wrote its own checks, ran the whole thing on a made-up example, and found the bug mine could not. It had quietly deleted part of my plan and said everything was fine.

How it runs, *top to bottom*.



Whoever builds it does not audit it.

Eight failures *my own checks could not see.*

The reviewer's first run, then the same checks today.

FIRST RUN, WRITTEN COLD

INDEPENDENT AUDIT · FIRST RUN · WRITTEN COLD FROM THE FILES

Forty checks it wrote itself. Thirty-two pass. Eight fail.

THE EIGHT THAT FAILED

- ✗ a stale canvas silently deletes plan boxes added since the last render
- ✗ plan ids colliding with derived line ids force a false match in the diff
- ✗ delete every stage on the canvas, apply, re-render, crash
- ✗ the diff misses a dropped open question or a dropped why
- ✗ a one-line text edit the canvas does not re-wrap is missed
- ✗ shift-lines on a plan destroys it
- ✗ a plan with no name renders 'undefined' into the title
- ✗ a plan with a human box and no stages, crash

AFTER THE FIXES

THE SAME SUITE · AFTER THE FIXES · TODAY

Forty checks. Forty pass.

PIPELINE HAZARDS

- ✓ a stale canvas silently deletes plan boxes added since the last render
 - ✓ plan ids colliding with derived line ids must not force a false match
 - ✓ the diff notices when a plan's open question or why is dropped in the build
 - ✓ readback refuses deletions on a stale canvas, and writes a backup first
- suites 7 · tests 40 · pass 40 · fail 0

Green again. Which now means something, because the checks were not written by the person who built it.

The worst one never crashed. It quietly deleted part of my plan, kept no copy, and said everything was fine.

The brief, *so you can run it on yours.*

Six lines. They work on anything you have built and want checked.

THE GIVE · THE AUDIT BRIEF, SO YOU CAN RUN IT ON YOURS

How to get your work checked by something that did not build it

- A separate reviewer with no memory of the build, given only the files.
- It writes its checks before it reads any of your answers.
- It runs the whole thing start to finish on a made-up example.
- It ranks what it finds by what would hurt first, not by what is easiest to fix.
- It reports and never fixes. Fixes happen after, separately, then the checks run again.
- It says which checks it could not run, so you know what is still a guess.

Distilled from the brief I used. Whoever builds it does not audit it.

Green means more when someone else wrote the tests.